

AmhTTS: A Low-Resource Amharic Text-to-Speech System

Walegn Tewabe Sewunetie¹, Rahel Mekonen Tamiru², Israel D. Gebru²

¹AIMS Research & Innovation Centre, Kigali, Rwanda

²Ethiopian Artificial Intelligence Institute (EAI), Addis Abeba, Ethiopia

³Meta Codec Avatars Lab in Pittsburgh, USA

Correspondence: wsewunetie@aimsric.org

Abstract

The scarcity of high-quality, phonetically diverse datasets for low-resource languages like Amharic remains a significant barrier to high-fidelity speech synthesis. We present **AmhTTS**, the first open, natively evaluated Amharic text-to-speech system, developed by fine-tuning SpeechT5 on a curated 25-hour, single-speaker Amharic corpus paired with a HiFi-GAN vocoder. Unlike Meta MMS, the only other publicly available Amharic TTS system, which requires Latin script transliteration and reports no Amharic-specific evaluation, AmhTTS operates natively on Ethiopic script and is the first system to report human MOS scores on Amharic speech synthesis. To address the orthographic complexities of the Ethiopic script, we implement a three-stage text normalisation pipeline covering homophone unification, *Ge'ez* numeral expansion, and Amharic punctuation handling. System quality is assessed through a rigorous human evaluation via a purpose-built web platform, involving **28 native Amharic-speaking voluntary evaluators** who rated 35 linguistically motivated samples across four domains: phonetic coverage (ejectives and labialized consonants), gemination and prosody, numeral and date reading, and conversational fluency. Results yield an overall naturalness MOS of **3.71** and intelligibility MOS of **4.02** on a 5-point scale, demonstrating that the rich pre-trained representations of SpeechT5 transfer effectively to a low-resource language with only 25 hours of data a fraction of the 460 hours used in the original English fine-tuning. While the system achieves high phonetic accuracy (MOS: 4.22/4.36), numeral synthesis and gemination modeling remain areas for future improvement. Given the demonstrated community demand for accessibility tools, we publicly release our fine-tuned model weights, the AmhTTS corpus, and evaluation data to advance low-resource TTS research. Our web application is currently active, serving

users at multiple university disability centers in Ethiopia.

<https://huggingface.co/amhtts>.

1 Introduction

Amharic is the official working language of Ethiopia, spoken natively by approximately 57 million people and used as a lingua franca by over 25 million additional speakers (Ethnologue, 2023). Despite its scale, Amharic remains severely under-resourced in speech technology, with no standardized TTS evaluation benchmark and very few publicly available systems Sebsibe et al. (2004).

Amharic presents several unique challenges for TTS synthesis. The Ge'ez/Ethiopic script encodes consonants and vowels in syllabic units (*fidel*), where multiple characters share identical pronunciations for example, **ሀ**, **ሐ**, and **ኸ** all realise /ha/ creating ambiguity during training. **Gemination** is phonemically contrastive yet difficult for neural models to learn reliably, while **ejective** and **labialized consonants** are typologically absent from English, limiting naive transfer learning. Additionally, Ethiopic numerals (**፩**, **፪**, ..., **፳**) and punctuation (**፡**, **፣**, **፥**) require dedicated normalisation pipelines absent from general-purpose pre-trained models.

Recent neural TTS architectures from the end-to-end Tacotron (Wang et al., 2017) and Tacotron 2 (Shen et al., 2018), to the non-autoregressive FastSpeech 2 (Ren et al., 2021) and the vocoderless VITS (Kim et al., 2021) have advanced synthesis quality significantly, but remain focused on high-resource languages. Meta MMS (Pratap et al., 2023), the only other publicly available Amharic TTS, provides no Amharic-specific evaluation and requires Latin transliteration via uroman, limiting practical usability. SpeechT5 (Ao et al., 2022), pre-trained exclusively on English, offers rich cross-modal representations

that we show transfer effectively to Amharic with only 25 hours of data.

In this work, we introduce **AmhTTS**, the first open, natively evaluated Amharic TTS system, fine-tuned from SpeechT5 (Ao et al., 2022) on a curated 25-hour single-speaker corpus with a HiFi-GAN vocoder (Kong et al., 2020). A three-stage text normalisation pipeline handles homophone unification, numeral expansion, and punctuation filtering. Human evaluation by 28 native Amharic-speaking evaluators yields a naturalness MOS of **3.71** and intelligibility MOS of **4.02**. Our main contributions are:

- The first open, fully reproducible Amharic TTS system operating natively on Ethiopic script without transliteration, with human MOS evaluation by native speakers.
- A three-stage text normalisation pipeline addressing homophone unification, *Ge'ez* numeral expansion, and Amharic punctuation handling.
- A linguistically motivated human evaluation framework combining ITU-T P.800 MOS (ITU-T, 1996) with forced transcription (Wester et al., 2015) involving 28 native Amharic-speaking evaluators across four phonetic domains (ejectives, gemination, numerals, and conversational fluency).
- Public release of model weights, the 25-hour AmhTTS corpus, evaluation data, and a live demo system to support reproducibility and Afro-Asiatic NLP research.

The remainder of this paper is organised as follows. Section 2 reviews related work. Section 3 describes the dataset and normalisation pipeline. Section 4 presents the model and fine-tuning. Section 5 details evaluation and results. Section 6 discusses findings and community impact. Section 7 concludes the paper, and Section 8 outlines the limitations of this work.

2 Related Work

2.1 Amharic Speech Synthesis

Early Amharic TTS relied on rule-based, cepstral, and concatenative synthesis with limited prosodic control (Sebsibe et al., 2004). More recent efforts have adapted neural sequence-to-sequence architectures, including Tacotron 2 (Shen et al., 2018)

and FastSpeech 2 (Ren et al., 2021) variants trained on modest corpora, with reporting naturalness scores between 3.2 and 3.8 MOS. The only other publicly available Amharic TTS system, Meta MMS (Pratap et al., 2023), covers 1,107 languages using a VITS-based (Kim et al., 2021) architecture but provides no Amharic-specific MOS evaluation and requires Latin transliteration via uroman. Table 1 summarises the landscape of relevant TTS systems; AmhTTS is the only system to report human MOS scores natively evaluated on Amharic.

2.2 Neural TTS Architectures

End-to-end TTS has progressed from the seq2seq Tacotron (Wang et al., 2017) and Tacotron 2 (Shen et al., 2018), to non-autoregressive models such as FastSpeech 2 (Ren et al., 2021) with explicit pitch, energy, and duration predictors. VITS (Kim et al., 2021) eliminated the need for a separate vocoder via VAE and GAN-based training, and serves as the backbone of Meta MMS (Pratap et al., 2023). Zero-shot multilingual models such as YourTTS (Casanova et al., 2022) and XTTS-v2 (Casanova et al., 2024) support voice cloning across languages with minimal data, though neither supports Amharic. SpeechT5 (Ao et al., 2022) unified speech and text in a shared encoder-decoder framework pre-trained on large-scale English data. AmhTTS adopts SpeechT5 as its backbone, demonstrating for the first time its cross-lingual transferability to a low-resource Afro-Asiatic language.

2.3 TTS Evaluation and African Language NLP

The ITU-T P.800 MOS framework is the standard for TTS naturalness assessment (ITU-T, 1996), while forced transcription provides a more reliable intelligibility measure for low-resource languages (Wester et al., 2015). AmhTTS combines both. The Mozilla Common Voice project has begun collecting Amharic speech (Ardila et al., 2020), and Masakhane-Speech (Olatunji et al., 2023) has produced benchmarks for several African languages. However, community-scale Amharic TTS evaluation with native speakers remains absent from the literature, a gap this work addresses.

Table 1: Comparative metrics for TTS systems relevant to Amharic (2026). MOS scores marked with † are reported on English benchmarks in the original papers, included for architectural reference only. Meta MMS Amharic checkpoint (facebook/mms-tts-amh) is released as part of the MMS project (Pratap et al., 2023) but Amharic is not explicitly evaluated in the paper. N/A = not reported. N/I = Naturalness / Intelligibility.

Model	Year	Key Feature	MOS (N/I)	Data Size	Language(s)	Data	Model
Tacotron (Wang et al., 2017)	2017	First end-to-end seq2seq TTS from characters	3.82† / N/A	24.6 hrs	English	Closed	Open
Tacotron 2 (Shen et al., 2018)	2018	Mel-spectrogram prediction conditioned WaveNet vocoder	4.53† / N/A	~24 hrs	English	Closed	Open
FastSpeech 2 (Ren et al., 2021)	2021	Non-autoregressive; explicit pitch, energy & duration	3.83† / N/A	~24 hrs	English	Open	Open
VITS (Kim et al., 2021)	2021	End-to-end VAE+GAN; no separate vocoder needed	4.43† / N/A	~24 hrs	English	Open	Open
YourTTS (Casanova et al., 2022)	2022	Zero-shot multi-speaker multilingual; ~1 min audio	3.98† / N/A	~1k spkrs	Multilingual (3)	Open	Open
SpeechT5 (Ao et al., 2022)	2022	Unified encoder-decoder pre-training for speech & text	3.65† / N/A	460 hrs	English	Open	Open
AmhTTS (Ours)	2025	SpeechT5 fine-tuned on Amharic; first open natively evaluated system	3.71 / 4.02	25 hrs	Amharic	Open	Open

† Amharic checkpoint at facebook/mms-tts-amh; not evaluated in Pratap et al. (2023).

3 Dataset and Text Normalisation

3.1 Training Corpus

AmhTTS is built on a high-fidelity, single-speaker corpus of **25 hours** of speech recorded by a female native Amharic speaker, comprising approximately **11,000 sentences**. The corpus contains **254 unique Ethiopic characters** with **747,546** total occurrences, ranging from a minimum count of 1 (፩, ፪, ፫) to a maximum of 43,622 (፳፻፲፱). Table 2 summarises the corpus statistics. Figure 1 visualises the character frequency distribution.

Table 2: AmhTTS training corpus statistics.

Property	Value
No. of sentences	11,000
Max sentence length	23 sec
Min sentence length	5 sec
Avg sentence length	≈16 sec
Total duration	25 hours
Sampling rate	16,000 Hz
Speaker	Female, native Amharic
Unique characters	254
Total characters	747,546

3.2 Text Normalisation Pipeline

Amharic orthography (Ethiopic Abugida, Unicode U+1200–U+137F) presents pre-processing challenges absent from general-purpose pre-trained models. We implement a three-stage pipeline.

Stage 1: Homophone Unification. Several Ge’ez characters share identical phonetic realisations, creating training noise. We map variants to

canonical forms: {ሀ, ሐ, ኀ, ኸ} → ሀ; {ሰ, ሠ} → ሰ; {ጸ, ፀ} → ጸ. Table 3 summarises key script properties.

Table 3: Amharic script properties affecting TTS.

Property	Details
Base consonants	33, each with 7 vowel orders (231 base forms)
Homophone groups	ሀ/ሐ/ኀ; ሰ/ሠ; ጸ/ፀ
Ge’ez numerals	፩=1, ፪=2, ፫=100, ፬=10,000
Key punctuation	፥ stop, ፣ comma, ፤ semi, ፡ question, ፡ word sep.

Stage 2: Numeral Expansion. Both Arabic (e.g., 2024) and Ge’ez numerals (e.g., ፳፻፲፱) are expanded to full Amharic written-out equivalents. Without this step the model skips or mispronounces numeric tokens (G3: MOS 2.92/3.16).

Stage 3: Punctuation Filtering. Non-standard symbols are stripped, while essential markers ፥, ፣, ፤, ፡ are preserved and mapped to prosodic pause durations.

4 Model Architecture and Fine-Tuning

4.1 Base Model: SpeechT5

AmhTTS fine-tunes **SpeechT5** (Ao et al., 2022), a unified-modal encoder-decoder pre-trained on large-scale unlabeled English speech and text. The framework consists of a shared Transformer encoder-decoder backbone and six modal-specific pre/post-nets (Ao et al., 2022). For TTS, text tokens are mapped via a text pre-net; the encoder-




Figure 1: Frequency heatmap of the 254 Ge'ez characters in the training corpus. Each cell shows the character glyph and its raw count. Colour intensity is logarithmic (light blue = rare, deep navy = frequent). Characters are ordered left-to-right, top-to-bottom by decreasing frequency.

decoder models the sequence-to-sequence transformation; and the speech post-net refines predicted log Mel-filterbank features. **HiFi-GAN** (Kong et al., 2020) converts mel-spectrograms to waveforms. Speaker identity is conditioned via x-vector embeddings. The original SpeechT5 was pre-trained and evaluated exclusively on English; AmhTTS is among the first efforts to adapt it to a low-resource morphologically rich language.

4.2 Fine-Tuning Procedure

Fine-tuning adapts the SpeechT5 attention mechanism to Ethiopic syllabic units using the normalised corpus (Section 3.2). Training uses L1

loss on log Mel-filterbank features and binary cross-entropy (BCE) loss on the stop token, following the original SpeechT5 procedure (Ao et al., 2022). All experiments used a single **NVIDIA Tesla V100-SXM2** (32 GB GPU), fully utilised over **≈30.5 hours** of training via Transformers 4.38.1 (Wolf et al., 2020) and PyTorch 2.1.0. Table 4 summarises all hyperparameters.




Figure 2: Training and validation loss over 16,000 steps (Weights & Biases). *Left*: Training loss decreases from 0.5244 (step 1k) to 0.4487 (step 16k). *Right*: Validation loss drops from 0.4757 (step 1k) to a minimum of **0.4305** at step 11,000, then stabilises at 0.4319 at the final checkpoint.

Table 4: Fine-tuning hyperparameters and compute details.

Hyperparameter	Value
Base model	microsoft/speecht5_tts
Vocoder	HiFi-GAN
Learning rate	1e-5
Train / Effective batch	16 / 128
Gradient accum. steps	8
Optimizer	Adam ($\beta_1=0.9$, $\beta_2=0.999$)
LR scheduler	Linear + 500 warmup steps
Total steps / Epochs	16,000 / ≈ 262
Sampling rate	16,000 Hz
Seed	42
Hardware	1 $\times$ V100-SXM2 (32 GB)
Training runtime	≈ 30.5 hrs
Framework	Transformers 4.38.1 / PyTorch 2.1.0
Best eval loss	0.4305 (step 11,000)
Final eval loss	0.4319 (step 16,000)

4.3 Training Dynamics

Figure 2 shows the loss curves over 16,000 steps. Validation loss drops from 0.4757 (step 1,000) to a minimum of **0.4305** at step 11,000, then stabilises at 0.4319. The plateau after step 9,000 is consistent with the limited data regime typical of low-resource fine-tuning. The slight uptick in training loss at step 16,000 (0.4487 vs. 0.4428 at step 15,000) with stable validation loss indicates no significant overfitting, attributable to the learning rate decaying to $\approx 1e-6$. The released model uses the **final checkpoint** (step 16,000, val. loss 0.4319), consistent with standard practice. This loss is within the 0.40–0.48 range reported for similar multilingual encoder-decoder fine-tuning on 20–40 hour African language datasets (Elmadany et al., 2025; Bakoubolo et al., 2026).

5 Evaluation and Results

5.1 Test Dataset

We constructed 35 test sentences in four groups (Appendix B): **G1** phonetic coverage ejectives, labialized consonants, 6th-order vowels (S1–S10); **G2** gemination and prosody minimal pairs (S11–S20); **G3** numbers, dates, and punctuation (S21–S25); **G4** conversational and domain-specific speech (S26–S35).

5.2 Results

Quality was assessed via ITU-T P.800 MOS (ITU-T, 1996) with forced transcription for intelligibility (Wester et al., 2015), collecting naturalness and intelligibility ratings (1–5) from **28 evaluators** (26 native, 2 fluent non-native Amharic speakers) and **30 community users**. As shown in Table 5, AmhTTS achieves naturalness MOS of **3.71** and intelligibility MOS of **4.02**, exceeding SpeechT5’s English baseline (Ao et al., 2022) (3.65 on 460 hrs) with only 25 hours of Amharic data. All 28 evaluators expressed willingness to use the system (“Yes”: 14; “Maybe”: 14).

Table 5: Overall MOS ($N=28$, 35 sentences).

Metric	Score	Ref. Data
Naturalness (per-sent.)	3.71	—
Intelligibility (per-sent.)	4.02	—
Naturalness. (self-reported)	3.61	—
Intelligibility. (self-reported)	3.93	—
SpeechT5 (Ao et al., 2022)	3.65	460 hrs
Tacotron (Wang et al., 2017)	3.82	24.6 hrs

6 Discussion

6.1 Comparison with Related Systems

As shown in Table 1, AmhTTS achieves naturalness MOS 3.71 and intelligibility MOS 4.02, exceeding the English SpeechT5 baseline (Ao et al., 2022) (3.65 on 460 hours) with only 25 hours of Amharic data. This demonstrates effective cross-lingual transfer of pre-trained encoder-decoder representations consistent with the broader transfer learning literature (Ren et al., 2021; Kim et al., 2021). AmhTTS is the first system to report native Amharic MOS and to operate directly on Ethiopic script without transliteration.

6.2 Results by Sentence Group

Table 6 reveals a clear hierarchy. G1 (phonetic coverage) achieved the highest scores (4.22/4.36), confirming solid ejective consonant handling. G4 (conversational) also scored strongly (3.73/4.18), reflecting good coverage of high-frequency vocabulary. G2 (gemination) shows a naturalness–intelligibility gap (3.58 vs. 3.96), suggesting evaluators recovered words even when gemination was imperfect, consistent with Sebsibe et al. (2004). G3 (numbers/dates) scored lowest (2.92/3.16), confirming numeral normalisation as the critical weakness.

Table 6: Mean MOS by sentence group.

Group	N	Nat.	Int.
G1: Phonetic Coverage	10	4.22	4.36
G2: Gemination/Prosody	10	3.58	3.96
G3: Numbers/Dates	5	2.92	3.16
G4: Conversational	10	3.73	4.18
Overall	35	3.71	4.02

6.3 Phonetic Error Analysis

Table 7 summarises phonetic error annotations across 980 observations. No errors were reported in 41.9% of cases. Gemination (17.9%) and ejective errors (16.3%) were the leading types, while number errors (4.9%) concentrated in G3. The lowest-scoring sentence, S21 (ዛሬ ሐምሌ ፩ ቀን ፳፻፲፯ ዓ.ም ነው።, Ge’ez numeral date), received naturalness 1.39 and intelligibility 1.36. Conversely, S5 (ኢትዮጵያ የብዙ ቋንቋዎች መገኛ ናት) achieved the highest scores (4.82/4.86), confirming strong synthesis on high-frequency vocabulary.

Table 7: Phonetic error distribution ($N=980$).

Error Type	Count	%
No errors detected	410	41.8
Other issues	175	17.9
Gemination unclear	175	17.9
Ejective sounds unclear	160	16.3
Number pronunciation unclear	60	6.1
Total	980	100.0

6.4 Failure Analysis

Gemination is not encoded at the character level in Ethiopic script; the model must infer consonant duration from context a task for which SpeechT5, pre-trained on English, has no prior. Non-autoregressive models with explicit duration predictors (Ren et al., 2021) offer a natural solution. **Numeral synthesis** failures stem from incomplete Ge’ez numeral handling in the normalisation pipeline, as evidenced by S21. Additional evaluator feedback identified: (1) male voice variants; (2) context-dependent homophone disambiguation; (3) improved prosody at punctuation boundaries; and (4) code-switched Amharic–Latin text support.

7 Conclusion

We presented **AmhTTS**, the first open, natively evaluated Amharic TTS system, fine-tuned from SpeechT5 (Ao et al., 2022) on a 25-hour single-speaker corpus using a three-stage Ethiopic normalisation pipeline. Human evaluation by 28 native speakers yielded a naturalness MOS of **3.71** and an intelligibility MOS of **4.02**, exceeding the English SpeechT5 baseline despite using a fraction of the training data. Current limitations include single-speaker coverage, suboptimal gemination (G2: 3.58/3.96), and weak numeral synthesis (G3: 2.92/3.16), attributable in part to the English-only SpeechT5 pre-training, which introduces systematic bias against ejective consonants and gemination features absent from its pre-training distribution. Future work will incorporate an explicit duration module, comprehensive numeral normalisation, and code-switched Amharic–Latin text support. The system has been publicly released and is actively used by visually impaired students, demonstrating its practical value for accessibility and digital inclusion in Amharic-speaking communities.

8 Limitations

Linguistic and Speaker Diversity. AmhTTS is trained on a single female speaker (Addis Ababa accent), limiting generalisation to regional varieties such as Gojjam and Wollo Amharic. Gemination handling remains suboptimal (G2: 3.58/3.96), and numeral synthesis requires a dedicated component (G3: 2.92/3.16). While our web application is currently active across multiple university disability centers in Ethiopia, the evaluator and initial user pool skew toward urban, educated academic environments, which may not fully capture rural user demographics or varying literacy levels upon wider nationwide deployment.

Transfer Learning and Infrastructure Constraints. SpeechT5 (Ao et al., 2022) is pre-trained exclusively on English, introducing a systematic bias against ejective consonants and gemination features absent from the pre-training distribution. This architectural foundation may structurally limit fine-tuning performance independently of data quantity. Furthermore, because our tool is currently deployed as a centralized web application, its utility is tightly coupled with internet stability. Optimization for lightweight, offline mobile-edge deployment remains a critical upcoming challenge to support users in low-connectivity areas.

9 Ethical Considerations

We acknowledge the dual-use risk of TTS technology, including the potential for audio deepfakes or malicious voice cloning. However, the primary motivation of this work is to improve accessibility for visually impaired Amharic speakers and promote digital inclusion in Ethiopian higher education. To address these risks transparently, the model weights, the AmhTTS corpus, and evaluation data are released under an open licence to advance low-resource, accountable research. No personally identifiable information was collected from users or evaluators; university center participants engaged with the application voluntarily with informed consent.

References

Junyi Ao, Rui Wang, Long Zhou, Chengyi Wang, Shuo Ren, Yu Wu, Shujie Liu, Tom Ko, Qing Li, Yu Zhang, and 1 others. 2022. [SpeechT5: Unified-modal encoder-decoder pre-training for spoken language](#)

[processing](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5723–5738.

Rosana Ardila and 1 others. 2020. Common voice: A massively-multilingual speech corpus. In *Proceedings of LREC*.

Justin E. Bakoubolo, Catherine Nana Nyaah Essuman, Messan Agbobli, Ahoefa Kansier, Eli Jean Doumeyan, Julie Pato, Notou Your Timibe, Emile KOGBEDJI Agossou, and Guedela Bakouya. 2026. [Eyaa-tom 26, yodi-mantissa and lom bench: A community benchmark for TTS in local languages](#). In *Proceedings of the 7th Workshop on African Natural Language Processing (AfricanNLP 2026)*, pages 264–270. Association for Computational Linguistics.

Edresson Casanova, Kelly Davis, Eren Gölge, Görkem Göknaç, Iulian Gulea, Logan Hart, Aya Aljafari, Joshua Meyer, Reuben Morais, Samuel Olayemi, and Julian Weber. 2024. [XTTS: a massively multilingual zero-shot text-to-speech model](#). *arXiv preprint arXiv:2406.04904*.

Edresson Casanova, Julian Weber, Christopher D Shulby, Arnaldo Candido Junior, Eren Gölge, and Moacir A Ponti. 2022. [YourTTS: Towards zero-shot multi-speaker TTS and zero-shot voice conversion for everyone](#). In *Proceedings of the 39th International Conference on Machine Learning (ICML)*, pages 2709–2720.

AbdelRahman Elmadany, Sang Yun Kwon, Hawau Olamide Toyin, Muhammad Abdul-Mageed, and 1 others. 2025. [Voice of a continent: Mapping Africa’s speech technology frontier](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics.

Ethnologue. 2023. [Amharic](#).

ITU-T. 1996. Methods for subjective determination of transmission quality. Technical Report P.800, International Telecommunication Union.

Jaehyeon Kim, Jungil Kong, and Juhee Son. 2021. [Conditional variational autoencoder with adversarial learning for end-to-end text-to-speech](#). In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, pages 5530–5540.

Jungil Kong, Jaehyeon Kim, and Jaekyoung Bae. 2020. HiFi-GAN: Generative adversarial networks for efficient and high fidelity speech synthesis. In *Advances in Neural Information Processing Systems*, volume 33, pages 17022–17033.

Tobi Olatunji, Tejumade Afonja, Aditya Yadavalli, Chris Chinenye Emezue, Sahib Singh, Bonaventure F. P. Dossou, Joanne Osuchukwu, Salomey Osei, Atnafu Lambebo Tonja, Naome Etori, and Clinton Mbataku. 2023. [AfriSpeech-200: Pan-African accented speech dataset for clinical and general domain](#)

ASR. Transactions of the Association for Computational Linguistics, 11:1669–1685.

Vineel Pratap, Andros Tjandra, Bowen Shi, Paden Tomasello, Arun Babu, Saranya Kundu, Ali Elkahky, Zhaoheng Ni, Apoorv Vyas, Maryam Fazel-Zarandi, and 1 others. 2023. [Scaling speech technology to 1,000+ languages](#). In *arXiv preprint arXiv:2305.13516*.

Yi Ren, Chenxu Hu, Xu Tan, Tao Qin, Sheng Zhao, Zhou Zhao, and Tie-Yan Liu. 2021. [FastSpeech 2: Fast and high-quality end-to-end text to speech](#). In *International Conference on Learning Representations (ICLR)*.

Yared Sebsibe, Solomon Kebede, and 1 others. 2004. Development of an Amharic text-to-speech system using cepstral method. In *Proceedings of the Workshop on Computational Approaches to Arabic Script-based Languages*, pages 83–89.

Jonathan Shen, Ruoming Pang, Ron J Weiss, Mike Schuster, Navdeep Jaitly, Zongheng Yang, Zhifeng Chen, Yu Zhang, Yuxuan Wang, RJ Skerry-Ryan, and 1 others. 2018. Natural TTS synthesis by conditioning WaveNet on mel spectrogram predictions. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 4779–4783.

Yuxuan Wang, RJ Skerry-Ryan, Daisy Stanton, Yonghui Wu, Ron J Weiss, Navdeep Jaitly, Zongheng Yang, Ying Xiao, Zhifeng Chen, Samy Bengio, and 1 others. 2017. [Tacotron: A fully end-to-end text-to-speech synthesis model](#). *arXiv preprint arXiv:1703.10135*.

Mirjam Wester, Cassia Valentini-Botinhao, and Gustavo Eje Henter. 2015. [Are we using enough listeners? No! — An empirically-supported critique of Interspeech 2014 TTS evaluations](#). In *Proceedings of Interspeech*, pages 3476–3480.

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, and 1 others. 2020. Transformers: State-of-the-art natural language processing. In *Proceedings of EMNLP: System Demonstrations*, pages 38–45.

A Appendix

B Test Dataset: 35 Amharic Sentences with MOS Scores

Table 8: Full test dataset with per-sentence MOS scores. (Nat. = Naturalness; Int. = Intelligibility)

S	Sentence	Nat.	Int.
Group 1: Phonetic Coverage (Ejective & Labialized)			
1	ጳውሎስ በሊባኖስ ተራራ ላይ ቆመ። (ጳ, ቆ)	4.18	4.61
2	ጨረቃ በጨለማ ብርሃን ትሰጣለች። (ጨ, ቃ)	4.71	4.71
3	ጥንቃቄ ለጤና በጣም አስፈላጊ ነው። (ጥ, ጤ)	4.57	4.75
4	ፀሐይ ጠዋት በምስራቅ በኩል ትወጣለች። (ፀ, ሐ)	4.57	4.82
5	ኢትዮጵያ የብዙ ቋንቋዎች መገኛ ናት። (ጵ)	4.82	4.86
6	የጭቃ ሹም በፈረስ ላይ ተቀመጠ። (ጭ, ሹ)	4.21	3.96
7	ቅዳሜና እሁድ ገበያ በጣም ይሞላል። (ቅ, ሁ)	4.36	4.86
8	መጽሐፍ ማንበብ እውቀትን ይጨምራል። (ጽ, ሐ)	4.71	4.75
9	ውሃ ጠጪ ብሎ አዘዛት። (Labialization)	3.39	3.75
10	ፍየሉ ቁጥቋጦውን በላው። (ቁ, ቋ)	2.64	2.54
Group 2: Gemination & Prosody			
11	ገና አልመጣም። (ገና = “yet”)	2.79	2.86
12	ገና በዓልን በደስታ አሳለፍን። (ገና = “Christmas”)	3.25	3.32
13	አባባ ደብዳቤ ጻፈ። (Plosive ‘b’)	2.50	2.82
14	እሱ ትልቅ ሰው ነው። (‘S’ sounds)	4.54	4.71
15	መመላለስ ሰልችቶኛል። (Repeated)	3.43	3.86
16	እሷ በጣም ጎበዝ ተማሪ ናት። (Feminine pitch)	4.43	4.82
17	ወንድሜ ትናንት መጣ። (Gemination ‘n’)	3.68	4.29
18	ይህ ሻይ በጣም ይፋጃል። (Affricates)	3.11	3.86
19	አንበሳው በጫካው ውስጥ ይጮኻል። (ኸ)	3.82	4.39
20	ሰማዩ ዛሬ በጣም ሰማያዊ ነው። (Vowels)	4.25	4.64
Group 3: Numbers, Dates & Punctuation			
21	ዛሬ ሐምሌ ፭ ቀን ፳፻፲፯ ዓ.ም ነው። (Ge’ez numerals)	1.39	1.36
22	በክፍሉ ውስጥ 12 ተማሪዎች አሉ። (Arabic numerals)	4.61	4.82
23	ሰዓቱ አራት ከሠላሳ ሆኗል። (Time format)	3.00	3.14
24	ዋጋው 500 ብር ከ60 ሳንቲም ነው። (Currency)	2.96	3.25
25	ስልክ ቁጥሩ 0911-22-33-44 ነው። (Phone numbers)	2.64	3.21
Group 4: Conversational & Domain-Specific			
26	እንደምን አደሩ? ዛሬ እንዴት ነዎት? (Formal)	3.14	3.64
27	የአየር ሁኔታ ትንበያው ዝናብ እንደሚዘንብ ተናግሯል። (News)	4.00	4.11
28	ኢንተርኔት ዓለምን ወደ አንዲት መንደር ቀይሯል። (Tech)	4.39	4.68
29	ጤና ይስጥልኝ! ስሜ ዋለልኝ ይባላል። (Intro)	2.86	3.57
30	ቡና መጠጣት ትወዳለህ ወይስ ሻይ? (Question)	2.93	3.68
31	የአፍሪካ አንድነት ድርጅት ዋና መሥሪያ ቤት አዲስ አበባ ነው። (Political)	4.21	4.57
32	ባቡሩ በአስር ደቂቃ ውስጥ ይደርሳል። (Transport)	3.89	4.50
33	ምሳ ለመብላት የት እንሂድ? (Casual)	3.18	3.75
34	ትምህርት ለትውልድ ዕድገት ወሳኝ ነው። (Educational)	4.29	4.61
35	ኳስ መጫወት የአካል ብቃት እንቅስቃሴ ነው። (Sports)	4.36	4.71