

Assessing Large Language Models as Automated Writing Evaluation and Feedback Tools: A Human–AI Agreement Study on L2 English Writing

Kowshik Saha¹, Moyen Uddin², Md Tanvirul Islam Turad³, and Mizanur Rahman⁴

^{1,2,3}Department of Computer Science and Engineering

⁴Department of Humanities

Bangladesh University of Engineering and Technology

¹kowshiksahakabya@gmail.com, ²2005075@cse.buet.ac.bd, ³2005011@cse.buet.ac.bd,

⁴mizanurrahman@hum.buet.ac.bd

Abstract

Large language models (LLMs) are often suggested as tools for grading student writing and giving feedback, but it is unclear how far we can trust them on the separate skills that make up good writing. We compare two LLM graders (Gemini and Llama) with eight professional English teachers on 20 paragraphs written in second-language English by Bengali-speaking undergraduates in Bangladesh, each graded by two teachers, scoring content, organization, vocabulary, and grammar separately as well as the total. We then run a second stage in which teachers read each model’s written explanation, judge whether they agree with it, and rate the feedback. The LLMs match the teachers on the total score, content, and organization, but are unreliable on vocabulary and grammar, where the two models do not even agree with each other. We also show that one low agreement score is misleading: it comes from graders giving very similar marks (score clustering), not from real disagreement. In the second stage, teachers found the AI feedback acceptable and useful most of the time ($\geq 77.5\%$), with very few made-up errors (2.5%), and preferred Gemini’s stricter grammar score to their own 20% of the time. LLMs can help grade writing, but a teacher should audit their output with extra care on vocabulary and grammar.

1 Introduction

Teachers spend much of their time grading student writing and writing feedback on it. Large language models (LLMs) can now do both, which makes them attractive as low-cost tools (Mizumoto and Eguchi, 2023; Stahl et al., 2024). But a single overall score hides a question: a model may judge overall quality well and still be unreliable on the specific skills, such as grammar or word choice, that useful feedback has to address.

We therefore ask three questions. **(RQ1)** On which writing skills do LLM graders agree with

expert teachers, and on which do they disagree? **(RQ2)** When teachers see a model’s score together with its written explanation, are the disagreements acceptable, or are they real mistakes? **(RQ3)** Is the feedback the models produce useful, appropriately worded, and free of invented errors?

Each paragraph was graded independently by two teachers, which lets us measure how much teachers naturally differ. We test Gemini 3 Flash Preview (Google DeepMind, 2025) and Llama 3.3 70B Instruct (Meta AI, 2024), the latter through Groq (Groq, 2024). Alongside the gold-standard Quadratic Weighted Kappa (Cohen, 1968) we use Gwet’s AC2, which stays reliable under score clustering (Gwet, 2008). Our contributions are: (i) a skill-by-skill picture of when LLM grading matches teachers (total score, content, organization) and when it does not (vocabulary, grammar); (ii) evidence that a chance-corrected measure is needed, so that low agreement caused only by score clustering is not mistaken for real disagreement; and (iii) an audit showing that most score gaps are acceptable and that the feedback is useful, appropriately worded, and free of invented errors.

2 Related Work

Early automated essay scoring (AES) systems predicted scores from hand-picked text features (Attali and Burstein, 2006); later ones used neural networks (Uto, 2021); see Ke and Ng (2019) and Ramesh and Sanampudi (2022) for surveys. Recent studies ask an LLM to grade essays with no training examples or only a few (Mizumoto and Eguchi, 2023; Mansour et al., 2024), to judge how well a text holds together (Naismith et al., 2023), to place writers on a proficiency scale (Yancey et al., 2023), and to score and give feedback at once (Stahl et al., 2024), following the wider idea of an LLM as a judge (Zheng et al., 2023). Most report one overall agreement number. We instead report agreement

for each writing skill, pair QWK with Gwet’s AC2, which is less sensitive to score clustering (Gwet, 2008; Feinstein and Cicchetti, 1990), and ask teachers to check the models’ explanations and feedback, since feedback is what helps students improve (Hattie and Timperley, 2007).

3 Methodology

3.1 Corpus and Rubric

We use 20 student paragraphs written on the prompt “Fake News on Social Media.” The writers are first-year undergraduates at a public university in Bangladesh. Their first language is Bengali, so this is second-language English writing. We have no independent measure of their English proficiency. Each paragraph is graded on a rubric with four separate skills: Content & Relevance (max 3.0), Organization & Cohesion (max 2.0), Lexis & Academic Register (max 2.5; *vocabulary* for short), and Grammar & Mechanics (max 2.5), adding up to a total of 10.0 in steps of 0.25 (Appendix B). Both LLM graders get the same prompt (Appendix A) and no example gradings: Gemini 3 Flash Preview (gemini-3-flash-preview) through the Gemini API (Google DeepMind, 2025), and Llama 3.3 70B Instruct (Meta AI, 2024) through Groq’s llama-3.3-70b-versatile endpoint at temperature 0.3 (Groq, 2024). We also fix the order of the answer (Appendix C): a score and a written reason for each skill, then the total, then the feedback. This makes the model look at each skill carefully before it commits to a single overall score.

3.2 How Paragraphs Were Assigned to Teachers

We split the 20 paragraphs among 8 professional English teachers at two Bangladeshi public universities (4 men, 4 women) using a distributed cyclic incomplete block design, so that no teacher’s workload is heavy or uneven. One of the eight is an author of this paper and wrote the rubric; all eight were equally informed about the study, graded five paragraphs, and were blind to the AI’s output in Phase 1. The LLM graders ran first, so we used Gemini’s total to sort the corpus into low (below 5.0), average (5.0 to 6.0), and high (above 6.0) quality, giving 4, 12, and 4 paragraphs. This binning only balances workload across blocks and plays no part in the agreement analysis. The bins went into eight overlapping blocks, so that every paragraph is graded by *exactly two* teachers and

every teacher grades one low, three average, and one high quality paragraph. Phase 1 therefore gives 40 human ratings, two per paragraph: two lists of 20 for the human–human comparison, and all 40 paired against each model’s score, repeated for both raters. We keep both teacher scores rather than averaging them, so the spread between teachers is not lost.

3.3 Two-Phase Protocol

In **Phase 1 (blind scoring)**, teachers grade their paragraphs. A teacher has to finish Phase 1 before Phase 2 opens, so nothing the AI produced is visible to them while they are scoring. In **Phase 2 (audit)**, each teacher is shown the AI’s score together with its written explanation. Where the two scores differ on a skill, the teacher picks one of three answers: (1) the AI’s score is better, and its explanation pointed out something they had missed (we call this being *persuaded*); (2) their own score is better, and the AI’s explanation is unconvincing; or (3) the scores differ but both are reasonable. Teachers also rate on a 1 to 5 scale how easy the feedback is to act on and how appropriate its tone is, and report any error the AI flagged that is not actually there.

3.4 How We Measure Agreement

Quadratic Weighted Kappa (QWK; Cohen, 1968) is the standard measure in essay scoring and counts large score gaps as worse than small ones. Gwet’s AC2 with quadratic weights (Gwet, 2008) measures the same thing but is robust to score clustering. Pearson r and Spearman ρ show whether two graders move up and down together, even if one is consistently stricter. Adjacent Agreement (AA) is the share of paragraphs scored within 0.5 points for a single skill and within 1.0 point for the total, and Mean Absolute Error (MAE) is the average size of the gap. Because 20 paragraphs is a small sample, we also report 95% confidence intervals for QWK and AC2 from 1,000 bootstrap resamples. We resample paragraphs rather than rows, since the 40 rows are 20 paragraphs rated twice and the two rows for one paragraph are not independent.

4 Results and Analysis

4.1 Phase 1: The LLMs Stay Near the Middle

Figure 1 shows both LLMs keeping their scores close to the middle of the scale while teachers use the full range. On the total score, human marks are

widely spread (mean 5.975, SD 1.713), but Gemini (SD 0.939) and Llama (SD 0.743) stay in a narrow band. This is the worst on vocabulary and grammar.

4.2 Phase 1: Where the LLMs Do Well and Where They Fail

Table 1 gives every agreement measure, with confidence intervals, for each skill and pair of graders. On the *total* score, *content*, and *organization*, a human and an LLM agree about as much as two humans do. On the total score, human–AI QWK is 0.462 for Gemini and 0.422 for Llama against 0.490 human–human, and the human–AI intervals $[0.18, 0.62]$ and $[0.19, 0.57]$ overlap substantially with the human–human interval $[-0.09, 0.85]$, so these data cannot separate them. Human–AI AC2 is higher (0.825 and 0.821 against 0.738), and content behaves the same way. The scores then drop on *vocabulary* (0.225 and 0.212 against 0.433) and *grammar* (0.194 and 0.007 against 0.324).

4.3 Why One Measure Is Not Enough

No single measure tells the whole story, and three cases show why.

The *kappa paradox* (Feinstein and Cicchetti, 1990) is clearest on organization: the teacher pair has almost no QWK (0.052) but a reasonable AC2 (0.574). The teachers do agree closely on 80% of the paragraphs, and QWK is low only because of score clustering, not because they really disagree. The confidence intervals point the same way. For the two teachers, the QWK interval dips below zero on all five skills, so on the QWK scale we cannot rule out that they are only agreeing by chance. Their AC2 interval stays above zero on all five, so on the AC2 scale we can. One interval on 20 paragraphs would prove little, but the same split appearing on every skill is a sign that QWK is the shakier of the two measures at this sample size.

A model can also exhibit score clustering, as Figure 1 shows for vocabulary and grammar. Gemini’s vocabulary score is almost flat, and because it sits near the mark teachers usually give, it does well on the measures that only look at distance: against Llama, 82.5% versus 77.5% adjacent agreement and MAE 0.344 versus 0.556. But its Pearson r of 0.318 and Spearman ρ of 0.327 fall below the 0.435 and 0.375 between teachers, so it is barely telling the paragraphs apart.

Grammar then splits the two models in opposite ways. Gemini’s QWK (0.194) and AC2 (0.607) trail the teacher values (0.324 and 0.756), but its

r of 0.294 and ρ of 0.311 are close to their 0.338 and 0.285, so it sorts the paragraphs much as they do and simply marks harder. Llama is the reverse: an AC2 of 0.812 looks healthy, while a QWK of 0.007, an r of 0.010, and a ρ of 0.022 indicate strong score clustering around the typical teacher-assigned score.

4.4 Is a Teacher Still Needed?

The Gemini vs. Llama row of each block in Table 1 shows how much the two LLMs agree with each other. On the total score, content, and organization they agree *more* with each other than any pair involving a human does (total QWK 0.826, AC2 0.894). Agreement between models from different developers is a good sign, though it does not guarantee either is right. On vocabulary and grammar even this support is gone. The two models’ vocabulary QWK is 0.000, with an interval of $[-0.05, 0.03]$, the tightest in the study and centred on no agreement at all, and their grammar AC2 interval $[-0.16, 0.65]$ contains zero. On these two skills we cannot show that the models agree with each other better than chance.

4.5 Phase 2: What the Teachers Said

The audit tells us which Phase 1 gaps are real mistakes and which are only numbers behaving oddly (Tables 2 and 3). The acceptability rate, the share of judgments where the AI score matched the teacher, was accepted as better, or was called equally reasonable, is high everywhere (77.5% to 92.5%), so most score gaps are ones a teacher can live with rather than ones that would mislead a student. The most interesting result is on grammar: teachers said Gemini’s stricter score was better than their own 20% of the time, so its poor Phase 1 numbers partly reflect sound strictness rather than noise. The written feedback also scored well: mean actionability 3.90 for Gemini and 3.80 for Llama, mean tone 3.83 and 3.80, and an invented error in only 2.5% of cases for both models.

5 Conclusion

Given only the rubric and no worked examples, an LLM agrees with a teacher about as closely as two teachers agree with each other on *total quality*, *content*, and *organization*. It is not reliable on *vocabulary* and *grammar*, where the two models do not even agree with each other. A chance-corrected measure is needed, or a skill where everyone gives similar marks looks like a failure when it is not.

Score Distributions by Rater and Rubric Factor




Figure 1: Score spread by grader and rubric skill (box plots, individual scores as dots). Both LLMs stay near the middle compared to the teachers, most strongly on Lexis and Grammar.

Skill	Grader Pair	QWK [95% CI]	AC2 [95% CI]	r	ρ	AA (%)	MAE
Total	Human vs. Human	0.490 [−0.09, 0.85]	0.738 [0.34, 0.92]	0.502	0.426	70.0	1.150
	Human vs. Gemini	0.462 [0.18, 0.62]	0.825 [0.69, 0.87]	0.582	0.544	57.5	1.200
	Human vs. Llama	0.422 [0.19, 0.57]	0.821 [0.68, 0.88]	0.622	0.587	62.5	1.175
	Gemini vs. Llama	0.826 [0.70, 0.91]	0.894 [0.83, 0.95]	0.849	0.863	100.0	0.375
Content & Rel.	Human vs. Human	0.420 [−0.06, 0.78]	0.808 [0.50, 0.94]	0.444	0.408	80.0	0.350
	Human vs. Gemini	0.484 [0.11, 0.73]	0.878 [0.75, 0.94]	0.501	0.521	85.0	0.300
	Human vs. Llama	0.412 [0.11, 0.64]	0.864 [0.75, 0.92]	0.485	0.478	87.5	0.350
	Gemini vs. Llama	0.657 [0.30, 0.85]	0.904 [0.78, 0.96]	0.727	0.674	100.0	0.200
Org. & Coh.	Human vs. Human	0.052 [−0.34, 0.47]	0.574 [0.12, 0.84]	0.052	0.093	80.0	0.438
	Human vs. Gemini	0.312 [−0.02, 0.58]	0.840 [0.73, 0.92]	0.353	0.343	90.0	0.269
	Human vs. Llama	0.368 [0.04, 0.63]	0.821 [0.68, 0.92]	0.407	0.380	92.5	0.269
	Gemini vs. Llama	0.583 [0.20, 0.82]	0.951 [0.62, 0.98]	0.586	0.651	100.0	0.125
Lexis & Reg.	Human vs. Human	0.433 [−0.15, 0.84]	0.792 [0.47, 0.96]	0.435	0.375	90.0	0.338
	Human vs. Gemini	0.225 [−0.02, 0.48]	0.871 [0.72, 0.94]	0.318	0.327	82.5	0.344
	Human vs. Llama	0.212 [0.03, 0.41]	0.674 [0.48, 0.78]	0.514	0.506	77.5	0.556
	Gemini vs. Llama	0.000 [−0.05, 0.03]	0.570 [0.00, 0.72]	0.000	0.000	75.0	0.550
Grammar & Mech.	Human vs. Human	0.324 [−0.18, 0.75]	0.756 [0.41, 0.92]	0.338	0.285	80.0	0.400
	Human vs. Gemini	0.194 [−0.02, 0.40]	0.607 [0.29, 0.76]	0.294	0.311	55.0	0.575
	Human vs. Llama	0.007 [−0.26, 0.29]	0.812 [0.53, 0.92]	0.010	0.022	77.5	0.425
	Gemini vs. Llama	0.093 [−0.04, 0.24]	0.303 [−0.16, 0.65]	0.301	0.297	55.0	0.650

Table 1: Phase 1 agreement measures for every skill and pair of graders. Confidence intervals come from 1,000 bootstrap resamples drawn at the paragraph level. AA = Adjacent Agreement, MAE = Mean Absolute Error.

Model	Skill	Persu. (%)	Accept. (%)
Gemini	Content & Rel.	5.0	87.5
Llama	Content & Rel.	2.5	87.5
Gemini	Org. & Coh.	2.5	85.0
Llama	Org. & Coh.	7.5	92.5
Gemini	Lexis & Reg.	5.0	90.0
Llama	Lexis & Reg.	10.0	82.5
Gemini	Grammar & Mech.	20.0	85.0
Llama	Grammar & Mech.	10.0	77.5

Table 2: Phase 2 persuasion and acceptability rates for each skill.

The audit shows that many apparent failures are differences a teacher accepts or cases where the AI was right, with few invented errors and feedback students can act on. We therefore suggest a teacher in the loop: let the model draft scores and feedback

Model	Action.	Tone	Halluc. (%)
Gemini	3.90	3.83	2.5
Llama	3.80	3.80	2.5

Table 3: Phase 2 mean rating for how easy the feedback is to act on (1–5), for its tone (1–5), and the share of judgements where the model reported an error that was not there (hallucination).

on all four skills, but keep a teacher auditing everything before the student sees it. Vocabulary and grammar need the most care, and our results indicate what to look for: score clustering may signal that the model has failed to distinguish a paragraph from the rest, so the teacher should check whether its score is appropriately differentiated rather than merely plausible on its own.

Limitations

Our study is small: 20 paragraphs, one writing prompt, two LLMs, and eight teachers. This is a resource constraint rather than a design choice. API costs add up quickly, and the human side is heavier still: a teacher scores every paragraph on four skills, then reads and audits the AI's score, reasoning, and feedback for each. A bigger study would mean either more paragraphs per teacher, which brings back the fatigue the block design avoids, or more teachers, which we could not afford. What we offer instead is the design, which carries over unchanged to a larger corpus and panel. The results are preliminary, so the agreement numbers are a rough guide rather than a settled answer.

All the paragraphs argue a point on one social science topic and are second-language English written by Bengali speakers, so our findings may not hold for other kinds of writing, other first languages, or first-language English writers. This matters most for grammar, where a writer's first language shapes which errors appear. We also have no independent measure of the writers' proficiency, so we cannot separate the effect of ability from the effect of first language. Few-shot examples or fine-tuning could change the vocabulary and grammar results.

The audit is done by the same teachers who gave the Phase 1 scores, so they may be attached to their earlier marks. One of the eight teachers is an author of this paper and wrote the rubric. That teacher did not take part in the study's design or in the analysis, did not know which agreement measures would be computed, and was informed about the study on the same terms as the other seven. The teacher appears in 5 of the 20 human-human pairs and contributes 5 of the 40 audit judgments per skill. Authorship is still an interest the other seven do not have, and it matters most in Phase 2, where the three response options differ visibly in how favourable they are even without knowledge of the metrics. At 5 of 40 judgments, however, no acceptability rate could shift by more than 12.5 points, and the observed rates range from 77.5% to 92.5%. We did not exclude these ratings, since dropping a quarter of the human-human pairs would leave the benchmark estimated on 15 paragraphs. A rubric's author may also apply it more consistently than other raters, which could raise the human-human agreement we compare the models against, though the same rater also contributes to the human-AI comparison.

With only two ratings per paragraph, the agree-

ment estimates are uncertain: the intervals in [Table 1](#) are wide, and most differences between grader pairs are not resolved at this sample size. Finally, we test one version of each model, and results may change as models are updated.

Ethics Statement

The 20 paragraphs are pre-existing coursework from past assessments, not written for this study. A public university in Bangladesh provided them and authorised their research use, on the condition that it is not named. Individual student consent was not obtained. Before any analysis, names, student identifiers, and institutional details were removed. No paragraph is reproduced; all reported results are aggregated. No student's grade or academic standing was affected. The AI-generated scores and feedback were never shown to students or used in any assessment decision.

Eight professional English teachers took part, four men and four women. They were informed of the study's purpose beforehand, took part voluntarily, and could withdraw at any time. They received no financial compensation. They are not individually identifiable from the results. Formal ethics review was not sought. We relied on the safeguards described above: de-identification, non-reproduction, no effect on grades, and voluntary participation.

Only de-identified paragraph text was sent to the model APIs. Llama 3.3 70B was accessed through Groq, which does not use API inputs and outputs for model training. Gemini was accessed through the free tier of the Google Gemini API. Google's terms permit free-tier inputs and outputs to be used to improve Google's products, including human review; paid tiers exclude this. This is a shortcoming: the paid tier is the appropriate choice for student-authored text, and we recommend it for future work. No personally identifying information was transmitted, and the topic is general public policy; individual risk is minimal.

The training data of both commercial models is not publicly documented. Any biases they carry regarding writer background, dialect, or first language may reach their scores. A single-prompt, single-language corpus cannot detect such effects. This is a further reason for the teacher-in-the-loop use we recommend: a human auditor is the only safeguard against biases we cannot measure.

References

- Yigal Attali and Jill Burstein. 2006. [Automated essay scoring with e-rater® v.2](#). *Journal of Technology, Learning, and Assessment*, 4(3).
- Jacob Cohen. 1968. [Weighted kappa: Nominal scale agreement provision for scaled disagreement or partial credit](#). *Psychological Bulletin*, 70(4):213–220.
- Alvan R. Feinstein and Domenic V. Cicchetti. 1990. [High agreement but low kappa: I. the problems of two paradoxes](#). *Journal of Clinical Epidemiology*, 43(6):543–549.
- Google DeepMind. 2025. [Gemini 3 flash preview](#). Gemini API documentation, Google AI for Developers. Model code: gemini-3-flash-preview. Accessed 2026.
- Groq. 2024. [Llama-3.3-70b-versatile model documentation](#). GroqCloud documentation. Endpoint llama-3.3-70b-versatile, served with TruePoint Numerics quantization. Accessed 2026.
- Kilem Li Gwet. 2008. [Computing inter-rater reliability and its variance in the presence of high agreement](#). *British Journal of Mathematical and Statistical Psychology*, 61(1):29–48.
- John Hattie and Helen Timperley. 2007. [The power of feedback](#). *Review of Educational Research*, 77(1):81–112.
- Zixuan Ke and Vincent Ng. 2019. [Automated essay scoring: A survey of the state of the art](#). In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence (IJCAI)*, pages 6300–6308.
- Watheq Ahmad Mansour, Salam Albatarni, Sohaila Eltanbouly, and Tamer Elsayed. 2024. [Can large language models automatically score proficiency of written essays?](#) In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 2777–2786, Torino, Italia. ELRA and ICCL.
- Meta AI. 2024. [Llama 3.3 70b instruct model card](#). Meta Llama Models repository. Accessed 2026.
- Atsushi Mizumoto and Masaki Eguchi. 2023. [Exploring the potential of using an ai language model for automated essay scoring](#). *Research Methods in Applied Linguistics*, 2(2):100050.
- Ben Naismith, Phoebe Mulcaire, and Jill Burstein. 2023. [Automated evaluation of written discourse coherence using GPT-4](#). In *Proceedings of the 18th Workshop on Innovative Use of NLP for Building Educational Applications (BEA)*, pages 394–403.
- Dadi Ramesh and Suresh Kumar Sanampudi. 2022. [An automated essay scoring systems: a systematic literature review](#). *Artificial Intelligence Review*, 55(3):2495–2527.
- Maja Stahl, Leon Biermann, Andreas Nehring, and Henning Wachsmuth. 2024. [Exploring LLM prompting strategies for joint essay scoring and feedback generation](#). In *Proceedings of the 19th Workshop on Innovative Use of NLP for Building Educational Applications (BEA 2024)*, pages 283–298, Mexico City, Mexico. Association for Computational Linguistics.
- Masaki Uto. 2021. [A review of deep-neural automated essay scoring models](#). *Behaviormetrika*, 48(2):459–484.
- Kevin P. Yancey, Geoffrey Laflair, Anthony Verardi, and Jill Burstein. 2023. [Rating short L2 essays on the CEFR scale with GPT-4](#). In *Proceedings of the 18th Workshop on Innovative Use of NLP for Building Educational Applications (BEA)*, pages 576–584.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. [Judging LLM-as-a-judge with MT-bench and chatbot arena](#). In *Advances in Neural Information Processing Systems 36: Datasets and Benchmarks Track*, volume 36.

A Grading Prompt

Both LLM graders were given the same prompt, shown in [Figure 2](#), with no example gradings. The placeholder {{Grading Standard}} was replaced by the full rubric in [Appendix B](#), {{Output Format}} by the template in [Appendix C](#), and {{Student Submission Text}} by the paragraph being graded.

B Grading Rubric

The rubric was written by one of the eight teachers, who is an author of this paper. [Table 4](#) gives it in full: the four skills, the highest score each one can earn, what the grader should look at, and what a paragraph has to do to earn each score. The total score is the sum of the four skills (max 10.0) and is given in steps of 0.25. This whole rubric was inserted into the prompt in place of {{Grading Standard}}, and the teachers graded from the same rubric.

C Output Format

[Figure 3](#) gives the template that replaced {{Output Format}} in the prompt. It fixes the order of the answer: a score and a reason for each skill, then the total, then the two kinds of feedback, so that the model looks at each skill carefully before it commits to a single overall score.

You are an expert English writing evaluator. Grade the student submission strictly based on the rubric provided.

Target Audience & Tone:

The student is a first-year undergraduate student (a fresher). Keep your feedback encouraging, direct, and highly logical. When explaining structural issues, frame them pragmatically: focus on how clear paragraph organization builds a cohesive, easy-to-follow argument. Avoid overly complex literary jargon; instead, emphasize how grammatical errors impede clear communication and how precise vocabulary strengthens the overall clarity and academic impact of their writing.

Your goal is to provide a clear score breakdown, followed by specific Corrective Feedback and Improvement Feedback.

Feedback Definitions:

- * **Corrective Feedback**: Identify specific errors in grammar, mechanics, and blatant structural failures. Quote the exact error from the text, explain logically *why* it is incorrect, and provide the corrected version.
- * **Improvement Feedback**: Identify areas where the writing is acceptable but could be elevated to hit the next tier of the rubric. Focus on vocabulary enhancement, deeper analysis, and better transitions. Suggest specific ways to rewrite or expand these sections.

Assignment Context

Title: Paragraph on Fake News on Social Media
Topic: Social Science

Grading Guidelines

{{Grading Standard}}

Maximum Score : 10.0

Output Format

{{Output Format}}

Student Submission

{{Student Submission Text}}

Figure 2: The zero-shot grading prompt given to both LLM graders.

Skill (Max)	Score	What earns this score
Content & Relevance (3.0)	<i>Focus</i>	Relevance, depth of argument, and specific evidence.
	3.0	Exceptional. Accurately defines fake news within the social media context (algorithms, virality) and provides a sophisticated impact analysis with a real-world example.
	2.0	Competent. Addresses the prompt well but may be slightly generic or miss the “social media” specific mechanisms (e.g. treats it like general rumors).
	1.0	Emerging. Vague or surface-level; fails to provide an example or meaningful impact analysis.
Organization & Cohesion (2.0)	<i>Focus</i>	Paragraph unity, topic sentences, and transitions.
	2.0	Strong “Point-Evidence-Explanation” (PEE) structure. Clear topic sentence and seamless transitions between the definition and the consequences.
	1.0	Functional. Has a beginning, middle, and end, but the logical flow is “clunky” or the topic sentence is weak.
	0.5	Disorganized. Ideas are scattered; lacks a cohesive argument or logical progression.
Lexis & Academic Register (2.5)	<i>Focus</i>	Word choice, technical terminology, and formal tone.
	2.5	Precise. Correctly uses technical terms (disinformation, algorithmic bias, echo chambers). Maintains a formal, objective stance (avoids “I think” or “it’s scary”).
	1.5	Appropriate. Standard vocabulary. Understandable but relies on common words rather than academic or field-specific terminology.
	0.5	Informal. Uses conversational language, slang, or overly emotional/subjective phrasing.
Grammar & Mechanics (2.5)	<i>Focus</i>	Syntax, punctuation, spelling, and capitalization.
	2.5	Highly Accurate. Little to no errors in tense, subject-verb agreement, or sentence structure. Proper use of complex sentences.
	1.5	Minor Issues. 1–3 slips that do not impede meaning (e.g. a missing comma or a singular/plural slip).
	0.5	Pervasive Errors. Significant issues with syntax or “Mechanics” (spelling/punctuation) that make the paragraph difficult to read.

Table 4: The full grading rubric. The listed values are the anchor points; graders could give any score in between in steps of 0.25.

Use proper spacing and indentation. Format your response exactly as follows:

Score Breakdown:

1. Content & Relevance: [Score] / [Max Score]
<<indentation>>Reason: [Reason for the score]
2. Organization & Cohesion: [Score] / [Max Score]
<<indentation>>Reason: [Reason for the score]
3. Lexis & Academic Register: [Score] / [Max Score]
<<indentation>>Reason: [Reason for the score]
4. Grammar & Mechanics: [Score] / [Max Score]
<<indentation>>Reason: [Reason for the score]

Overall Score: [Total Score] / 10.0

Corrective Feedback:

- <<Numbered list in the following format>>
- <<Number>> [Quote the specific error from the text]
<<indentation>> Issue: [Briefly explain the grammatical or structural error]
<<indentation>> Correction: [Provide the fixed version]

Improvement Feedback:

- <<Numbered list in the following format>>
- <<Number>> [Quote the specific sentence/phrase that could be improved]
<<indentation>> Target Area: [Name the rubric category being addressed, e.g., Lexis & Academic Register]
<<indentation>> Suggestion: [Explain how to elevate this specific text to reach a higher rubric tier, offering a concrete example. Suggest a better alternative for the sentence.]

Figure 3: The output format given to both LLM graders. Per-skill scores and reasons come before the total score, which comes before the feedback.